

Homework 2 - Generalized Hough Transform

Theory

< Insert your answers here >

Programming

Find object in an image using a template:



```

In [1]: #!/usr/bin/env python3
        # -*- coding: utf-8 -*-
        import cv2
        import utils
        import numpy as np
        from matplotlib import pyplot as plt
        from sklearn.metrics.pairwise import euclidean_distances

def nonMaxSupression(img, d=5):
    """
    Given an image set all values to 0 that are not
    the maximum in its (2d+1,2d+1)-window

    Parameters
    -----
    img : ndarray
        an image
    d : int
        for each pixels consider the surrounding (2d+1,2d+1)-window

    Returns
    -----
    result : ndarray

    """

    rows, cols = img.shape
    result = np.zeros((rows, cols))

    # iterate over pixels
    # iterate over (2d+1,2d+1) neighborhood window
    for row in range(rows):
        for col in range(cols):
            left, right = max(0, col - d), min(cols - 1, col + d + 1)
            top, bottom = max(0, row - d), min(rows - 1, row + d + 1)

            # supress non-maxima to 0
            # store results in new array
            m = np.max(img[top:bottom, left:right])
            result[row, col] = m if img[row, col] == m else 0

```

```
return result
```

```
def calcBinaryMask(img, thresh=0.3):
```

```
    """
```

```
    Compute the gradient of an image and compute a binary mask  
    based on the threshold. Corresponds to 0^B in the slides.
```

```
    Parameters
```

```
    -----
```

```
    img : ndarray
```

```
        an image
```

```
    thresh : float
```

```
        A threshold value. The default is 0.3.
```

```
    Returns
```

```
    -----
```

```
    binary : ndarray
```

```
        A binary image.
```

```
    """
```

```
    # -compute gradients
```

```
    abs_gradient = np.abs(utils.calcDirectionalGrad(img))
```

```
    # -threshold gradients
```

```
    result = np.where(abs_gradient > thresh *  
                      np.max(abs_gradient), 1.0, abs_gradient)
```

```
    # -return binary mask
```

```
    return result
```

```
def correlation(img, template):
```

```
    """
```

```
    Compute a correlation of gradients between an image and a template.
```

```
    Note:
```

```
    You should use the formula in the slides using the fourier transform.  
    Then you are guaranteed to succeed.
```

```
    However, you can also compute the correlation directly.
```

```
    The resulting image must have high positive values at positions
```

with high correlation.

Parameters

`img` : ndarray
 a grayscale image
`template` : ndarray
 a grayscale image of the template

Returns

ndarray
 an image containing the correlation between image and template gradients.
"""

```
rows, cols = template.shape
```

```
def extend_img(img, new_shape):  
    tmp = np.zeros(shape=new_shape, dtype=type(img))  
    tmp[:rows, :cols] = img  
    return tmp
```

```
# -dft of the image, already computed in GeneralizedHoughTransform(...)  
DFT_I = img
```

```
# -compute gradient of the template  
Oi = utils.calcDirectionalGrad(template)
```

```
# -copy template gradient into larger frame  
Oi = extend_img(Oi, new_shape=img.shape)  
Ob = extend_img(calcBinaryMask(template, 0.75), new_shape=img.shape)
```

```
# -apply a circular shift so the center of the original template is in the  
# upper left corner  
Oi = utils.circularShift(Oi, cols // 2, rows // 2)  
Ob = utils.circularShift(Ob, cols // 2, rows // 2)
```

```
# -normalize template  
normalization_factor = 0  
for row in range(rows):  
    for col in range(cols):  
        normalization_factor += np.abs(Oi[row, col])  
Oi /= normalization_factor
```

```

T = Oi * Ob

# -compute correlation
DFT_T = np.fft.fft2(T)
I_T = np.fft.ifft2(DFT_I * np.conj(DFT_T))
return np.abs(I_T.real)

def GeneralizedHoughTransform(img, template, angles, scales):
    """
    Compute the generalized hough transform. Given an image and a template.

    Parameters
    -----
    img : ndarray
        A query image
    template : ndarray
        a template image
    angles : list[float]
        A list of angles provided in degrees
    scales : list[float]
        A list of scaling factors

    Returns
    -----
    hough_table : list[(correlation, angle, scaling)]
        The resulting hough table is a list of tuples.
        Each tuple contains the correlation and the corresponding combination
        of angle and scaling factors of the template.

        Note the order of these values.
    """
    import itertools

    lst = []

    DFT_I = np.fft.fft2(utils.calcDirectionalGrad(img))

    # for every combination of angles and scales
    for angle, scale in itertools.product(angles, scales):

        # -distort template

```

```
distorted_template = utils.rotateAndScale(template, angle, scale)

# -compute the correlation
corr = correlation(DFT_I, distorted_template)

# -store results with parameters in a list
lst.append((corr, angle, scale))

return lst
```

Main Program

```
In [2]: # Load query image and template
query = cv2.imread("data/query.jpg", cv2.IMREAD_GRAYSCALE)
template = cv2.imread("data/template.jpg", cv2.IMREAD_GRAYSCALE)

# Visualize images
utils.show(query)
utils.show(template)



# Create search space and compute GHT
angles = np.linspace(0, 360, 36)
scales = np.linspace(0.9, 1.3, 10)
ght = GeneralizedHoughTransform(query, template, angles, scales)


# extract votes (correlation) and parameters
votes, thetas, s = zip(*ght)

# Visualize votes
votes = np.stack(votes).max(0)
plt.imshow(votes)
plt.show()


# nonMaxSuppression
votes = nonMaxSuppression(votes, 20)
plt.imshow(votes)
plt.show()

# Visualize n best matches
n = 10
coords = zip(
    *np.unravel_index(np.argpartition(votes, -n, axis=None)[-n:], votes.shape))
vis = np.stack(3*[query], 2)
for y, x in coords:
    print(x, y)
    vis = cv2.circle(vis, (x, y), 10, (255, 0, 0), 2)
utils.show(vis)
```





133 151
242 132
257 158
171 56
138 117
198 193
102 91
225 193
207 3
253 95



Test your implementation

```
In [3]: import utils
import cv2
import json
from matplotlib import pyplot as plt
import numpy as np
from sklearn.metrics.pairwise import euclidean_distances
```

```

In [4]: from sklearn.metrics.pairwise import euclidean_distances

def testGHT():
    query = cv2.imread("data/query.jpg", cv2.IMREAD_GRAYSCALE)
    template = cv2.imread("data/template.jpg", cv2.IMREAD_GRAYSCALE)

    angles = np.linspace(0, 360, 36)
    scales = np.linspace(0.9, 1.3, 10)
    ght = GeneralizedHoughTransform(query, template, angles, scales)

    votes, thetas, s = zip(*ght)
    votes = np.stack(votes).max(0)
    plt.imshow(votes)
    plt.show()

    #votes = correlation(query, template)
    votes = nonMaxSuppression(votes, 20)
    plt.imshow(votes)
    plt.show()

    n = 10
    coords = list(zip(*np.unravel_index(np.argpartition(votes, -n, axis=None)[-n:], votes.shape)))

    vis = np.stack(3*[query],2)
    for y,x in coords:
        vis = cv2.circle(vis,(x,y), 10, (255,0,0), 2)
    utils.show(vis)

    f = open("centroids.txt", "r")
    centroids = f.read()
    f.close()

    centroids = centroids.split("\n")[:-1]
    centroids = [centroid.split() for centroid in centroids]
    centroids = np.array([[int(centroid[0]),int(centroid[1])] for centroid in centroids])

    vis = np.stack(3*[query],2)
    for x,y in centroids:
        vis = cv2.circle(vis,(x,y), 10, (255,0,0), 2)
    utils.show(vis)

```

```
coords = np.array(coords)[:,:,:-1]



d = euclidean_distances(centroids, coords).min(1)



correct_detections = np.count_nonzero((d<10))

score = { "scores": {"Correct_Detections": correct_detections }}

print(json.dumps(score))

testGHT()
```





```
{"scores": {"Correct_Detections": 6}}
```